

# CS471: Introduction to Artificial Intelligence

## Assignment 4: Decision Tree Classification

In this assignment, you will implement the Decision tree classification method using **Scikit-learn**.

This is an individual assignment. Use Google Colab for your code, plots, and comments. When you finish editing, re-run all the cells to make sure they work.

---

In this assignment, you will work with the Iris dataset that was used in [R. A Fisher's paper](#). You can also find the dataset in the [UCI Machine Learning Repository](#). You can directly load the dataset using sklearn:

```
from sklearn.datasets import load_iris
data = load_iris()
```

Tasks:

1. Include a basic description of the data (what are the features and labels) (1 point)

Write in your own words of what the classification task is and why a decision tree is a reasonable model to try for this data. (1 point)

2. Split the data into training, validation, and testing sets. (1 point)
3. Fit a decision tree on the training dataset. (1 point)
4. Tune at least 2 hyperparameters in the decision tree model (<https://ken-hoffman.medium.com/decision-tree-hyperparameters-exp>)

[lained-49158ee1268e](#)) based on the performance on the validation set or using cross-validation. One hyperparameter has to be max\_depth and the other one is your choice.

Generate plot of hyperparameter values vs performance metric (plots are mandatory). (4 points)

5. Train the model using optimal hyperparameters (found in step 5) on the train + validation data. Test it on test data and generate a classification report (1 point)
6. Inspect the model by visualizing and interpreting the results (1 point)

Google Colab Link:

<https://colab.research.google.com/drive/1uwhbQPILedXzWCxIRXIM2ecPhDmZUeRy?usp=sharing>

GitHub Link:

<https://github.com/jacobalmon/CS-471/blob/main/Homework/Homework%2004/dtclass.py>

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split,
cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
```

These are the following modules used for implementing this assignment. Sklearn is mainly used to create the Machine Learning Model using a Decision Tree Classifier and Splitting Data. The matplotlib module was used to help plot the graphs for the validation dataset and the hyperparameters.

```

iris = load_iris() # Load Iris Dataset.
x = iris.data # Features.
y = iris.target # Labels.

# Split Data into Training, Validation, and Testing
Datasets.
x_train, x_test, y_train, y_test = train_test_split(x,
y, test_size=0.2, random_state=29)
x_train, x_valid, y_train, y_valid =
train_test_split(x_train, y_train, test_size=0.2,
random_state=29)

```

In this following code segment, we are loading the iris dataset and breaking the dataset into a training, validation, and testing datasets. We just define some random state, in our case its 29.

```

max_depth_vals = range(1, 15)
min_samples_leaf_vals = range(1, 11)

# Creating Lists for the Outcomes of the Hyper
Parameters.
max_depth_outcomes = []
min_samples_leaf_outcomes = []

# Evaluating Decision Tree Classifier for Max Depth.
for max_depth in max_depth_vals:
    dt_classifier =
DecisionTreeClassifier(max_depth=max_depth,
random_state=29)
    dt_classifier.fit(x_train, y_train)
    outcomes = dt_classifier.score(x_valid, y_valid)
    max_depth_outcomes.append(outcomes)

```

```

    # Evaluating Decision Tree Classifier for Min Samples
    Leaf.
    for min_samples_leaf in min_samples_leaf_vals:
        dt_classifier =
DecisionTreeClassifier(min_samples_leaf=min_samples_leaf,
random_state=29)
        dt_classifier.fit(x_train, y_train)
        outcomes = dt_classifier.score(x_valid, y_valid)
        min_samples_leaf_outcomes.append(outcomes)

```

In this following code segment, we are tuning the hyper parameters which are max depth and min samples leaf. We then evaluate the model with Decision Tree Classifier with each hyperparameter by storing the outcomes of each values of the ones we define. We keep the same random state of 29 as well.

```

plt.figure(figsize=(12,6))

# Plotting the Validation Accuracy vs. Max Depth.
plt.subplot(1,2,1)
plt.plot(max_depth_vals, max_depth_outcomes, marker='o')
plt.title('Validation Accuracy vs. Max Depth')
plt.xlabel('Max Depth')
plt.ylabel('Validation Accuracy')
plt.xticks(max_depth_vals)
plt.grid(True)

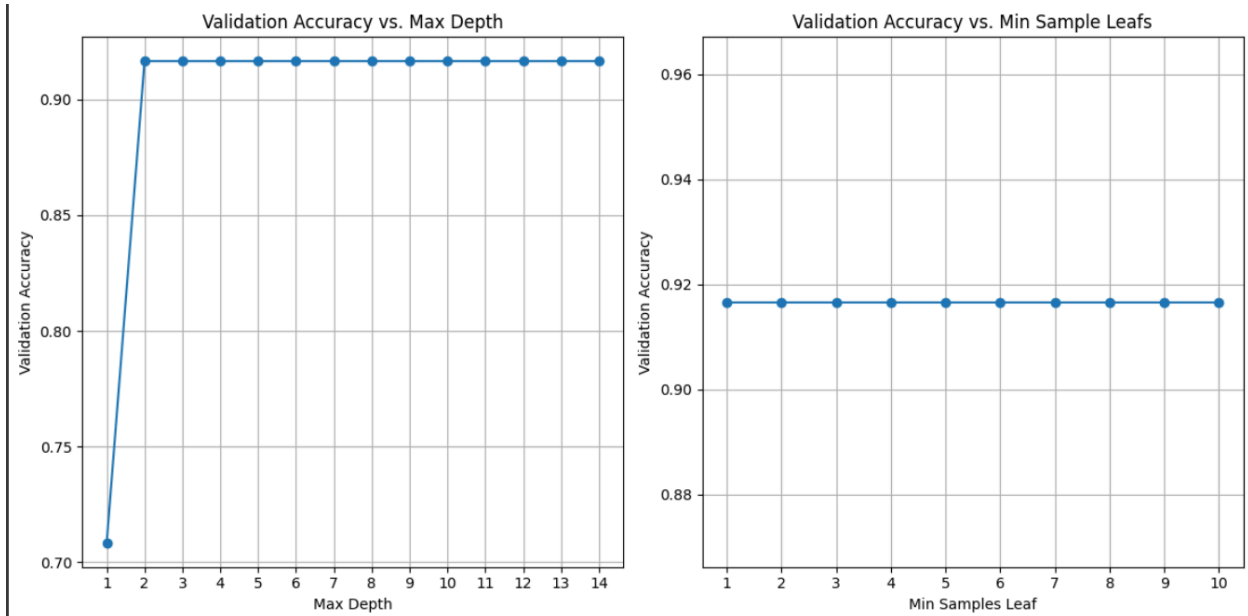
# Plotting the Validation Accuracy vs. Min Samples Leaf.
plt.subplot(1,2,2)
plt.plot(min_samples_leaf_vals, min_samples_leaf_outcomes, marker='o')
plt.title('Validation Accuracy vs. Min Samples Leafs')
plt.xlabel('Min Samples Leaf')
plt.ylabel('Validation Accuracy')
plt.xticks(min_samples_leaf_vals)
plt.grid(True)

# Display Plots onto the Screen.

```

```
plt.tight_layout()
plt.show()
```

In this following code segment we plot the two hyperparameters against the validation dataset. Our results are:



Initially the max\_depth would be low since there is only one datapoint to find the max\_depth, so it's not very accurate, but when we go up more nodes, we see the accuracy improves and stays constant above 90 percent. On the other hand the min sample leafs stay at a constant time below 92 percent in that plot.

```
# Find the Best Max Depth Outcome.
best_max_depth =
max_depth_vals[max_depth_outcomes.index(max(max_depth_outcomes))]
# Find the Best Min Samples Leaf Outcome.
best_min_samples_leaf =
min_samples_leaf_vals[min_samples_leaf_outcomes.index(max(min_samples_leaf_outcomes))]

# Fitting the Model with these new Best Max Outcome & Min Samples Leaf Outcome.
final_dt_classifier = DecisionTreeClassifier(max_depth=best_max_depth,
min_samples_leaf=best_min_samples_leaf, random_state=42)
final_dt_classifier.fit(x_train, y_train)

# Find the Test Accuracy from our Testing Data.
```

```
y_pred = final_dt_classifier.predict(x_test)
test_accuracy = accuracy_score(y_test, y_pred)

print(f'Test Dataset Accuracy: {test_accuracy * 100}')
```

After evaluating our ML model with the vadiation dataset, we need to pick the best one and we do this in the following code segment to find the best max depth and best min samples leaf. Lastly we pick the best model and test the model on our testing data which results the accuracy of:

```
Test Dataset Accuracy: 90.0
```